

Diffusion Generative Flow Samplers: Improving Learning Signals Through Partial Trajectory Optimization

Dinghuai Zhang*, Ricky T. Q. Chen, Cheng-Hao Liu, Aaron Courville &
Yoshua Bengio

Thomas Mousseau

October 28, 2025

Overview

1. Introduction

- 1.1 Problem Statement
- 1.2 Stochastic Optimal Control

2. Stochastic Optimal Control to GFlowNets

- 2.1 Credit Assignment Problem in SOC
- 2.2 GFlowNet

3. Diffusion Generative Flow Sampler Results

- 3.1 Gradients variance, Z estimation and mode coverage
- 3.2 Convergence Guarantees

4. Conclusion

- 4.1 Final Thoughts and Future Work
- 4.2 Q&A



Generative Modeling

Task

Sample from a complex (high-dimensional and multimodal) distribution D .

D can be given under the form of:

- A dataset of samples $\{x_i\}_{i=1}^N \sim D$ (e.g., images, text, audio)
- An unnormalized density $\mu(x)$ where D has density $\pi(x) \propto \mu(x)$ (e.g., energy-based models, physics/chemistry)



Sampling from Unnormalized Densities

Context. Sample from a D -dimensional target with unnormalized density $\mu(x)$ where $\mathbb{R}^D \rightarrow \mathbb{R}^+$.

$$\pi(x) = \frac{\mu(x)}{Z}, \quad Z = \int_{\mathbb{R}^D} \mu(x) dx \text{ (unknown)}.$$

We assume we can evaluate $\mu(x)$, but we have no samples from π and do not know Z .

Goal. We seek a *sampler* (similar to MCMC/VI) that produces calibrated samples and, ideally, estimates Z *without* any dataset from π .

Chemistry (molecule conformers).

Different 3D conformations have a formation energy from force-field terms (bonds, angles, dihedrals, nonbonded). A low energy means a higher probability of being sampled. A well-calibrated sampler is needed to draw conformers in proportion to these probabilities, which is important in binding-pose ranking, free-energy estimation and generating diverse realistic 3D conformers for screening.

Diffusion Generative Flow Samplers

Idea. We will *reframe sampling* from an unnormalized target $\pi(x) \propto \mu(x)$ as a *stochastic optimal control (SOC)* problem. This means learning a control function that steers a diffusion process so its *terminal marginal* matches π .

Why this helps.


- Gives a *path-space* training objective/metric: a KL on trajectories $\text{KL}(Q \parallel P)$ where P is the reference paths reweighted by $\mu(x_T)$.
- The *partition function* Z *cancels* inside this objective, so we can train using only μ (and optionally $\nabla \log \mu$).
- Lets us optimize *without samples from* π and still measure closeness to the true normalized endpoint.

Caveat (sets up DGFS). This path-KL places supervision *only at the terminal time* \Rightarrow poor *credit assignment* and high-variance gradients.

DGFS fix (preview). Inject *intermediate* learning signals via a GFlowNet-inspired *learned flow* and *subtrajectory balance*, enabling partial-trajectory training and more stable learning.

Diffusion Generative Flow Samplers v2

Minimizing the running and terminal costs between the optimal controlled process and the uncontrolled reference process.



$$\begin{aligned} & \underset{\tilde{\mathcal{P}}}{\text{minimize}} && D_{\text{KL}}[\tilde{\mathcal{P}} \parallel \mathcal{P}] = \int_{\Omega} \log \frac{d\tilde{\mathcal{P}}}{d\mathcal{P}} d\tilde{\mathcal{P}} \\ & \text{s.t.} && x(0) \sim \pi_0, \quad x(T) \sim \pi_T \end{aligned}$$

2025-10-28

Diffusion Generative Flow Samplers: Improving Learning Signals Through Partial Trajectory Optimization


- └ Introduction

- └ Problem Statement

- └ Diffusion Generative Flow Samplers v2

Diffusion Generative Flow Samplers v2

Minimizing the running and terminal costs between the optimal controlled process and the uncontrolled reference process.



refaire cette illustration avec le controlled, uncontrolled et target (optimal controlled) processus

Controlled and Reference Processes

Controlled forward transition (learned drift).

$$P_F(x_{n+1} | x_n) = \mathcal{N}(x_{n+1}; x_n + h f_\theta(x_n, n), h\sigma^2 I)$$

Controlled process.

$$Q(x_{0:N}) = p_0^{\text{ref}}(x_0) \prod_{n=0}^{N-1} P_F(x_{n+1} | x_n)$$

Uncontrolled/Reference forward transition (zero drift).

$$P_F^{\text{ref}}(x_{n+1} | x_n) = \mathcal{N}(x_{n+1}; x_n, h\sigma^2 I)$$

Uncontrolled/Reference process and marginals.

$$Q^{\text{ref}}(x_{0:N}) = p_0^{\text{ref}}(x_0) \prod_{n=0}^{N-1} P_F^{\text{ref}}(x_{n+1} | x_n),$$

$$p_n^{\text{ref}}(x) = \mathcal{N}(x; \mu_0, \Sigma_0 + nh\sigma^2 I), \text{ is closed form.}$$

Goal Learn f so that the terminal marginal $Q(x_N)$ matches $\pi(x) = \mu(x)/Z$ (no data, Z unknown)

Diffusion Generative Flow Samplers: Improving Learning Signals Through Partial Trajectory Optimization

- Introduction
- Stochastic Optimal Control
- Controlled and Reference Processes

Why closed form for uncontrolled but not controlled? The uncontrolled reference process has no drift, so each transition is a simple Gaussian convolution, leading to Gaussian marginals with closed-form means and variances (e.g., $p_n^{ref}(x) = N(x; \mu u_0, \text{cov}_0 + nh\sigma^2 I)$). The controlled process includes a learned drift $f(x_n, n)$, which makes transitions non-Gaussian and dependent on f , preventing a closed-form marginal expression without solving the integral numerically or via simulation.

je dois refaire cette slide, on ne voit meme comment x_n est sample (par rapport a la SDE) et aussi d'ou sort sigma mu et la closed form expliquer et c'est une convolution de gaussienne

Controlled and Reference Processes

Controlled forward transition (learned drift).

$$P_f(x_{n+1} | x_n) = N(x_{n+1}; a_n + h f_n(x_n, a_n), h\sigma^2 I)$$

Controlled process.

$$Q(x_0, x_N) = p_0^{ref}(x_0) \prod_{n=0}^{N-1} P_f(x_{n+1} | x_n)$$

Uncontrolled/Reference forward transition (zero drift).

$$P_0^{ref}(x_{n+1} | x_n) = N(x_{n+1}; x_n, h\sigma^2 I)$$

Uncontrolled/Reference process and marginals.

$$Q^{ref}(x_0, x_N) = p_0^{ref}(x_0) \prod_{n=0}^{N-1} P_0^{ref}(x_{n+1} | x_n).$$

$$p_n^{ref}(x) = N(x; \mu_0, \text{cov}_0 + nh\sigma^2 I), \text{ is closed form.}$$

Goal: show f so that the stochastic marginal $Q(x_n)$ matches $\pi(x)$ in L^2 (no data, π unknown).

Closed Form of Reference Marginal

Reference Process.

$$x_N = x_0 + \sqrt{h}\sigma \sum_{k=0}^{N-1} \varepsilon_k.$$

Distribution of Terminal State. Since the sum of independent Gaussians is Gaussian:

$$x_N \sim \mathcal{N}(m_0, \Sigma_0 + Nh\sigma^2 I).$$

Terminal Marginal Density.

$$p_N^{\text{ref}}(x) = \frac{1}{(2\pi)^{D/2} |\Sigma_0 + Nh\sigma^2 I|^{1/2}} \exp \left(-\frac{1}{2} (x - m_0)^\top (\Sigma_0 + Nh\sigma^2 I)^{-1} (x - m_0) \right).$$

Path space KL objective

Target path measure via terminal reweighting.

$$P(x_{0:N}) \propto Q^{\text{ref}}(x_{0:N-1}|x_N)\mu(x_N) \propto Q^{\text{ref}}(x_{0:N}) \frac{\mu(x_N)}{p_N^{\text{ref}}(x_N)} \implies P(x_N) \propto \mu(x_N).$$

KL decomposition.

$$\text{KL}(Q\|P) = \mathbb{E}_Q \left[\log \frac{Q}{P} \right] = \mathbb{E}_Q \left[\log \frac{Q}{Q^{\text{ref}}} \right] + \mathbb{E}_Q \left[\log \frac{p_N^{\text{ref}}(x_N)}{\mu(x_N)} \right] + \log Z.$$

Running control cost (Gaussian mean-shift) Girsanov theorem.

$$\mathbb{E}_Q \left[\log \frac{Q}{Q^{\text{ref}}} \right] = \mathbb{E}_Q \sum_{n=0}^{N-1} \frac{h}{2\sigma^2} \|f_\theta(x_n, n)\|^2.$$

Terminal cost.

$$\mathbb{E}_Q \left[\log \frac{p_N^{\text{ref}}(x_N)}{\mu(x_N)} \right] = \mathbb{E}_Q [\log p_N^{\text{ref}}(x_N) - \log \mu(x_N)]$$

Diffusion Generative Flow Samplers: Improving Learning Signals Through Partial Trajectory Optimization

- Introduction
 - Stochastic Optimal Control
 - Path space KL objective

Path space KL objective

Target path measure via terminal reweighting.

$$P\{\eta_{0:N}\} \propto Q^{\text{ref}}(\eta_{0:N-1}|\eta_0)\mu\{\eta_N\} \propto Q^{\text{ref}}(\eta_{0:N}) \frac{\mu\{\eta_N\}}{\rho\{\eta_N\}} \implies P\{\eta_N\} \propto \rho\{\eta_N\}.$$

KL decomposition.

$$\text{KL}(Q|P) = \mathbb{E}_Q \left[\log \frac{Q}{P} \right] = \mathbb{E}_Q \left[\log \frac{Q}{Q^{\text{ref}}} \right] + \mathbb{E}_Q \left[\log \frac{Q^{\text{ref}}(\eta_{0:N})}{\rho\{\eta_N\}} \right] + \log Z.$$

Running control cost (Gaussian mean-shift) Cramér theorem.

$$\mathbb{E}_Q \left[\log \frac{Q}{Q^{\text{ref}}} \right] = \mathbb{E}_Q \sum_{t=1}^{N-1} \frac{h}{2\pi^2} \|\dot{\eta}_t(\eta_{0:t-1})\|^2.$$

Terminal cost.

$$\mathbb{E}_Q \left[\log \frac{\mu\{\eta_N\}}{\rho\{\eta_N\}} \right] = \mathbb{E}_Q [\log \rho_{Q^{\text{ref}}}(\eta_N) - \log \rho\{\eta_N\}]$$

Je devrais avoir une slide complete a expliquer d'ou vient le foward target process P parce que c'est le coeur de mon algorithme

Target Path Measure via Importance Sampling

Importance Sampling Basics. Importance sampling is a technique to estimate expectations under a hard-to-sample target distribution P using samples from an easy-to-sample proposal distribution Q . The key formula is:

$$\mathbb{E}_{x \sim Q} [f(x) \cdot w(x)] = \mathbb{E}_{x \sim P} [f(x)],$$

where the *importance weight* $w(x) = \frac{P(x)}{Q(x)}$ corrects the samples from Q to behave as if they were from P .

Intuition. Q provides "biased" samples; $w(x)$ upweights samples that are likely under P and downweights those that aren't, effectively resampling from P without directly sampling it.

Application to Path Measures. In our case, the "samples" are entire trajectories $x_{0:N}$, and we want the path measure P to have the correct terminal marginal $\pi(x_N) \propto \mu(x_N)$. We use the reference path measure Q^{ref} (easy to sample, e.g., Gaussian paths) as the proposal. The reweighted path measure is:

$$P(x_{0:N}) \propto Q^{\text{ref}}(x_{0:N}) \cdot w(x_{0:N}),$$

where the weight $w(x_{0:N}) = \frac{\pi(x_N)}{P_N^{\text{ref}}(x_N)}$. This ensures $P(x_N) \propto \mu(x_N)$, as the weight depends only on the endpoint.

Credit Assignment Problem in SOC objective

SOC Discrete-time objective

$$\min_{f_\theta} \mathbb{E}_Q \left[\underbrace{\sum_{n=0}^{N-1} \frac{h}{2\sigma^2} \|f_\theta(x_n, n)\|^2}_{\text{Running cost}} + \underbrace{\log p_N^{\text{ref}}(x_N) - \log \mu(x_N)}_{\text{Terminal cost}} \right]$$

This objective is used on the seminar paper *Path Integral Sampler: Diffusion-based Sampling for Unnormalized Densities* Zhang et al. 2022a et al which presented the sampling from unnormalized densities as a stochastic optimal control problem.

Explanation. The SOC objective provides feedback signal only at the terminal step, making credit assignment difficult via backpropagation through time. This causes high-variance gradients, weak feedback for early actions, poor mode discovery, and inefficient optimization of the drift f without intermediate signals.

SOC as a GFlowNet

Comparison Table: GFlowNet vs. SOC Framework

Concept	GFlowNet	SOC
Forward Process	Trajectory sampling on DAG	Controlled diffusion path
Forward Transition Probability	$P_F(s' s)$	$P_F(x_{n+1} x_n) = \mathcal{N}(x_{n+1}; x_n + hf(x_n), h\sigma^2 I)$
Backward Transition Probability	$P_B(s s')$	$P_B^{\text{ref}}(x_n x_{n+1})$ (known)
Reward Function	$R(x)$ (unnormalized)	$\mu(x)$ (unnormalized density)
Terminal Marginal Distribution	$P_T(x) \propto R(x)$	$Q(x_N) \propto \mu(x_N)$
Flow State	Flow $F(s)$ at states	Learned flow $F_n(x)$

Insight. Since SOC can be viewed as a GFlowNet, we can apply GFlowNet tools (e.g., detailed balance loss, subtrajectory balance) to solve the credit assignment problem.

Diffusion Generative Flow Samplers: Improving Learning Signals Through Partial Trajectory Optimization

- └ Stochastic Optimal Control to GFlowNets
- └ Credit Assignment Problem in SOC
- └ SOC as a GFlowNet

SOC as a GFlowNet

Comparison Table: GFlowNet vs. SOC framework

Concept	GFlowNet	SOC
Forward Process	Trajectory sampling on DAG	Controlled diffusion path
Forward Transition Probability	$P_T(x' x)$	$P_T(x_{t+1} x_t) \propto \mathcal{N}(x_{t+1}; \kappa_t + \Delta F(x_t), \text{cov}^2 F)$
Backward Transition Probability	$P_B(x x')$	$P_B^{\text{ref}}(x_t x_{t+1})$ (known)
Reward Function	$R(x)$ (unnormalized)	$\mu(x)$ (unnormalized density)
Terminal Marginal Distribution	$P_T(x) \propto R(x)$	$Q(x_0) \propto \mu(x_0)$
Flow State	Flow $F(x)$ at states	Learned flow $F_T(x)$

Insight. Since SOC can be viewed as a GFlowNet, we can apply GFlowNet tools (e.g., detailed balance loss, subtrajectory balance) to solve the credit assignment problem.

make a mental note to comeback to professor hernandez comment on Tristan's Deleu presentation when someone asked: "The if GFlowNets are just Reinforcement Learning, then why keep doing research on this framework?" Professor Hernandez repoded by proving the felxibility of GFN, he said that it was en entery point between RL, Diffusion models, Energy-based models and in this case we prove that it can also be used in Stochastic Optimal Control problems.

Addressing Credit Assignment via Target Process

Target process necessitates terminal reweighting...

$$P(x_{0:N}) \propto Q^{\text{ref}}(x_{0:N}) \frac{\mu(x_N)}{p_N^{\text{ref}}(x_N)} \implies P(x_N) \propto \mu(x_N).$$

Could we write the target process differently to allow intermediate supervision?

Conditional Form. Since the reweighting only affects the terminal state, the joint can be written as:

$$P(x_{0:N}) = \pi(x_N) \prod_{n=0}^{N-1} Q_B^{\text{ref}}(x_n | x_{n+1}),$$

where $Q_B^{\text{ref}}(\cdot | \cdot)$ is the backward transition probability (derived from the target joint). This is tractable because Q_B^{ref} is known.

Diffusion Generative Flow Samplers: Improving Learning Signals Through Partial Trajectory Optimization

└ Stochastic Optimal Control to GFlowNets

└ GFlowNet

└ Addressing Credit Assignment via Target Process

conditional form is exactly like the RHS of the trajectory balance equation in gfn. Lire le paper de Nikolay sur Trajectory Balance pour bien comprendre cette equation et comment elle se relie a DGFS

Addressing Credit Assignment via Target Process

Target process necessitates terminal reweighting.

$$P(s_0, s) \propto Q^{\pi}(s_0, s) \frac{p(s_0)}{P^{\pi}(s_0)} \Rightarrow P(s_0) \propto p(s_0).$$

Could we write the target process differently to allow intermediate supervision?

Conditional Form. Since the reweighting only affects the terminal state, the joint can be written as:

$$P(s_0, s) = \pi(s_0) \prod_{i=0}^{N-1} Q^{\pi}(s_i, s_{i+1}),$$

where $Q^{\pi}(s_i, s_{i+1})$ is the backward transition probability (derived from the target joint). This is tractable because Q^{π} is known.

Rewriting the Target Process with Marginal

If We Had Access to $p_n(x_n)$ we could write the partial joint which would allow training on subtrajectories and thus have better credit assignment:

$$P(x_{0:n}) = p_n(x_n) \prod_{k=0}^{n-1} Q_B^{\text{ref}}(x_k | x_{k+1}),$$

But There's No Closed Form for $p_n(x_n)$. To calculate $p_n(x_n)$, we would need to compute the integral:

$$p_n(x_n) = \int \pi(x_N) \prod_n^{N-1} Q_B^{\text{ref}}(x_n | x_{n+1}) dx_{n+1:N}.$$

$$p_n(x_n) \propto \int_{\mathbb{R}^{(N-n)D}} \mu(x_N) \prod_{k=n}^{N-1} Q^{\text{ref}}(x_k | x_{k+1}), dx_{n+1:N}.$$

Why is $p_n(x_n)$ hard to compute? The integral represents summing over all possible future paths from x_n to the terminal x_N , weighting each by the terminal density $\pi(x_N)$. This high-dimensional integration (over $(N-n) \times D$ dimensions) has no analytical solution for general $\mu(x_N)$. Numerically approximating it via quadratures or Monte Carlo is computationally expensive, especially since training requires computing it repeatedly for subtrajectories as intermediate signals (replacing the single terminal signal), making it impractical.

Diffusion Generative Flow Samplers: Improving Learning Signals Through Partial Trajectory Optimization

└ Stochastic Optimal Control to GFlowNets

└ GFlowNet

└ Rewriting the Target Process with Marginal

ne pas oublier de parler des MC quadrature method pour approximer l'intégral mais c'est couteux en temps de calcul et pas scalable du tout

Rewriting the Target Process with Marginal

If We Had Access to $p_\theta(x_n)$ we could write the partial joint which would allow training on subtrajectories and thus have better credit assignment:

$$P(x_{0:n}) = p_\theta(x_n) \prod_{s=0}^{n-1} Q^{\theta^s}(x_s | x_{1:n}).$$

But There's No Closed Form for $p_\theta(x_n)$. To calculate $p_\theta(x_n)$, we would need to compute the integral:

$$p_\theta(x_n) = \int z(x_n) \prod_{s=0}^{n-1} Q^{\theta^s}(z_s | x_{1:n}) dz_{0:n-1}$$

$$p_\theta(x_n) \propto \int_{\mathbb{R}^N} p_\theta(x_n) \prod_{s=0}^{n-1} Q^{\theta^s}(z_s | x_{1:n}) dz_{0:n-1}$$

Why is $p_\theta(x_n)$ hard to compute? The integral represents summing over all possible future paths from x_n to the terminal x_0 , weighting each by the terminal density $p_\theta(x_0)$. This high-dimensional integration (over $(N-n) \times D$ dimensions) has no analytical solution for general $p_\theta(x_0)$. Numerically approximating it via quadrature or Monte Carlo is computationally expensive, especially since training requires computing it repeatedly for subtrajectories as intermediate signals (replacing the single terminal signal), making it impractical.

Trajectory Balance with Learned Flow

Proposed Amortized Approach. Train $F_n(\cdot; \phi)$, a neural network, to satisfy the following constraint for all partial trajectories $x_{n:N}$:

$$F_n(x_n; \phi) \prod_{k=n}^{N-1} P_F(x_{k+1}|x_k; \theta) = \mu(x_N) \prod_{k=n}^{N-1} P_B(x_k|x_{k+1}).$$

Details.

- P_F (forward policy) and F_n are parameterized by deep neural networks, with parameters θ and ϕ respectively.
- We can view F_n as an approximation to the unknown marginal $p_n(x_n)$ thus an amortized way to estimate the intractable integral.
- We use only $\mu(\cdot)$ (no Z), so the unknown normalization is absorbed into F_n .
- Very similar to a Value Function in RL, estimating future expected rewards (here, future density mass).

Detailed and SubTB Balance with Learned Flow

Detailed Balance from Trajectory Balance condition.

Comparing the constraint for n and $n + 1$ gives a formula independent of μ which now only involves local transitions (signal subtrajectories instead of full trajectories)

$$F_n(x_n; \theta) P_F(x_{n+1} | x_n; \theta) = F_{n+1}(x_{n+1}; \theta) P_B(x_n | x_{n+1}).$$

Subtrajectory Balance Loss for Partial Trajectories

$$\ell_{\text{SubTB}}(x_{m:n}; \theta, \phi) = \left(\log \frac{F_m(x_m; \phi) \prod_{k=m}^{n-1} P_F(x_{k+1} | x_k; \theta)}{F_n(x_n; \phi) \prod_{k=m}^{n-1} P_B(x_k | x_{k+1})} \right)^2$$

Using a mix of different subtrajectory lengths (m, n) allows better credit assignment thus more stable training of both the forward policy and the learned flow.

Diffusion Generative Flow Samplers: Improving Learning Signals Through Partial Trajectory Optimization

- Stochastic Optimal Control to GFlowNets

- GFlowNet

- Detailed and SubTB Balance with Learned Flow

a voir si je voudrais pas mettre une schema sur l'une des 3 derniers slides pour faire un recap de scrap la SOC objective, now we build DB loss with learned flow to solve credit assignment problem

Detailed and SubTB Balance with Learned Flow

Detailed Balance from Trajectory Balance condition

Comparing the constraint for n and $n+1$ gives a formula independent of ρ which now only involves local transitions (signal subtrajectories instead of full trajectories)

$$P_{\theta}(s_n; \theta) P_{\theta}(s_{n+1}; s_n; \theta) = P_{\theta}(s_{n+1}; \theta) P_{\theta}(s_n; s_{n+1})$$

Subtrajectory Balance Loss for Partial Trajectories

$$\ell_{\text{data}}(s_{n:n+1}; \theta, \phi) := \left(\log \frac{P_{\theta}(s_n; \phi) \prod_{i=0}^{n-1} P_{\theta}(s_{i+1}; s_i; \theta)}{P_{\phi}(s_n; \phi) \prod_{i=0}^{n-1} P_{\phi}(s_i; s_{i+1})} \right)^2$$

Using a mix of different subtrajectory lengths (n, s) allows better credit assignment than more stable training of both the forward policy and the learned flow.

Overall Training Objective

Recall SubTB Loss

$$\ell_{\text{SubTB}}(x_{m:n}; \theta, \phi) = \left(\log \frac{F_m(x_m; \phi) \prod_{k=m}^{n-1} P_F(x_{k+1} | x_k; \theta)}{F_n(x_n; \phi) \prod_{k=m}^{n-1} P_B(x_k | x_{k+1})} \right)^2$$

Diffusion Generative Flow Sampler (DGFS) Loss

$$L(\tau; \theta; \phi) = \frac{\sum_{0 \leq m < n \leq N} \lambda^{n-m} \ell_{\text{SubTB}}(x_{m:n})}{\sum_{0 \leq m < n \leq N} \lambda^{n-m}}, \quad \tau = (x_0, \dots, x_N)$$

This combines signals from all subtrajectory lengths, reducing variance and improving credit assignment

- τ : Full trajectory (x_0, \dots, x_N) .
- λ : Positive scalar weighting different subtrajectory lengths (e.g., shorter subtrajectories get higher weight if $\lambda < 1$).
- The numerator sums SubTB losses over all subtrajectories $x_{m:n}$, weighted by λ^{n-m} (length-based weighting).
- The denominator normalizes to average the losses..

Algorithm 1 DGFS Training

Require: $\mu(\cdot)$, $\bar{\sigma}$, N , λ , B , η

1: Init $\theta = (\theta_f, \phi)$

2: **repeat**

3: Sample trajectories:

4: **for** $b = 1$ to B **do**

5: $\tau^{(b)} = (x_0^{(b)}, \dots, x_N^{(b)})$ under $x_{n+1} = x_n + hf_{\theta_f}(x_n, n) + \sqrt{h}\bar{\sigma}\varepsilon_n$, $\varepsilon_n \sim \mathcal{N}(0, I)$

6: **end for**

7: Build subtrajectories: $\mathcal{S}(\tau^{(b)})$ of (m, n) with $0 \leq m < n \leq N$

8: Compute SubTB loss:

$$\mathcal{L}(\tau^{(b)}; \theta) = \sum_{(m,n) \in \mathcal{S}(\tau^{(b)})} \lambda^{n-m} \left[\log \frac{F_{\phi}(x_m) \prod_{l=m}^{n-1} P_F(x_{l+1} \mid x_l; \theta_f)}{F_{\phi}(x_n) \prod_{l=m}^{n-1} P_B^{\text{ref}}(x_l \mid x_{l+1})} \right]^2$$

9: $g \leftarrow \nabla_{\theta} \frac{1}{B} \sum_{b=1}^B \mathcal{L}(\tau^{(b)}; \theta)$

10: $\theta \leftarrow \text{Adam}(\theta, g, \eta)$

11: **until** convergence

Reduction of gradient variance




Figure: Gradient variance comparison between DGFS and PIS. DGFS shows significantly lower variance, leading to more stable training.




Figure: Drift network magnitude of DGFS and PIS.




Figure: Gradient variance of DGFS, PIS, and DGFS trained with only full paths

Accurate partition function Z estimation

How to Estimate Z . The partition function $Z = \int \mu(x) dx$ normalizes the density. Since we sample trajectories from the learned process Q , we estimate Z using importance sampling on the terminal states.

Log-partition function estimator:

$$\log \sum_{b=1}^B \exp(S(\tau^{(b)}; \theta)) - \log B \leq \log Z, \quad \tau^{(b)} \sim Q(\cdot; \theta).$$

Intuition (Simple).

- $S(\tau; \theta)$: Weight showing how well the trajectory matches the target (higher if it ends in high- μ regions).
- Sample many trajectories from Q , compute their weights, and average in log space.
- This gives a lower bound on $\log Z$ because Q approximates the target, providing an estimate without direct π samples.

Partition function estimation results

	MoG	FUNNEL	MANYWELL	VAE	Cox
SMC	0.289 ± 0.112	0.307 ± 0.076	22.36 ± 7.536	14.34 ± 2.604	99.61 ± 8.382
VI-NF	1.354 ± 0.473	0.272 ± 0.048	2.677 ± 0.016	6.961 ± 2.501	83.49 ± 2.434
CRAFT	0.348 ± 0.210	0.454 ± 0.485	0.987 ± 0.599	0.521 ± 0.239	13.79 ± 2.351
FAB w/ BUFFER ⁵	0.003 ± 0.0005	0.0022 ± 0.0005	0.032 ± 0.004	N/A	0.19 ± 0.04
PIS	0.036 ± 0.007	0.305 ± 0.013	1.391 ± 1.014	2.049 ± 2.826	11.28 ± 1.365
DDS	0.028 ± 0.013	0.416 ± 0.094	1.154 ± 0.626	1.740 ± 1.158	N/A ⁶
DGFS	0.019 ± 0.008	0.274 ± 0.014	0.904 ± 0.067	0.180 ± 0.083	8.974 ± 1.169

Figure: The lower the better, DGFS achieves the lowest bias in estimating the partition function across various benchmarks compared to PIS and DDS.

Convergence Guarantees

Convergence Guarantees for DGFS.

- Previous studies De Bortoli et al. 2022; Chen et al. 2022; Lee et al. 2022 show that the terminal sample distribution of a diffusion model converges to the target under mild assumptions if the control term is well learned. These apply to DGFS since proofs are independent of training method.
- Zhang et al. 2022b prove that a perfectly learned score corresponds to zero GFlowNet training loss E. Bengio, Jain, et al. 2021; E. Bengio, S. Bengio, et al. 2023, ensuring a well-trained DGFS accurately samples from the target distribution.

Diffusion Generative Flow Samplers: Improving Learning Signals Through Partial Trajectory Optimization

└ Diffusion Generative Flow Sampler Results

└ Convergence Guarantees

└ Convergence Guarantees

Le but de montrer cette slide de garanties de convergences est pour prouver que malgré les exemples simples MoG, Manywell, etc (doivent être simple pour avoir accès à la partition fonction Z , on peut garantir que DGFS va converger vers la bonne distribution cible pour des targets plus complexes.)

Convergence Guarantees

Convergence Guarantees for DGFS.

- Previous studies De Bortoli et al. 2022; Chen et al. 2022; Lee et al. 2022 show that the terminal sample distribution of a diffusion model converges to the target under mild assumptions if the control term is well learned. These apply to DGFS since proofs are independent of training method.
- Zhang et al. 2022b prove that a perfectly learned score corresponds to zero GFlowNet training loss. E. Bengio, Jain, et al. 2021; E. Bengio, S. Bengio, et al. 2023, ensuring a well-trained DGFS accurately samples from the target distribution.

Mode coverage results




Figure: Manywell plots. DGFS and DDS but not PIS recover all modes




Figure: MoG visualization of DGFS and other diffusion-based samplers shows that DGFS could capture the diverse modes well. The contours display the landscape of the target density

Strengths of DGFS

Strengths

- Works from unnormalized targets. Trains directly with $\mu(x)$ (and optionally $\nabla \log \mu$), no dataset or partition function Z needed.
- Well-calibrated endpoints. At convergence the terminal marginal $p_N(x)$ matches $\pi(x) \propto \mu(x)$; samples come with the right probabilities.
- Intermediate learning signals. Subtrajectory Balance (SubTB) provides supervision on partial paths $x_{m:n}$, solving the “terminal-only signal” issue of path-KL samplers (better temporal credit assignment, lower variance).
- Fixed, closed-form backward. By terminal tilting, the target’s backward kernel equals the reference bridge (Gaussian, analytic), giving stable, non-circular constraints.

Weaknesses of DGFS

Weaknesses / Limitations

- Reliance on Gaussian reference. Closed-form backward requires (near-)Gaussian reference with constant covariance; non-Gaussian or state-dependent noise breaks the neat math.
- Two nets to train. Must learn both the forward drift and the flow F_n ; more parameters, more tuning (time embeddings, loss weights).
- μ -query cost. Frequent evaluations of $\mu(x)$ (and possibly $\nabla \log \mu$) can be expensive; SubTB helps, but large-scale tasks need caching/replay/prioritization.
- Coverage still hard in high-D. Even with SubTB, exploring multi-modal, high-barrier landscapes can stall without tempering, momentum, or preconditioning.
- State-space must match target. DGFS assumes a fixed continuous state space \mathbb{R}^D ; discrete/structured spaces need bespoke references/backwards or relaxations.
- Hyperparameter sensitivity. Noise scale $\bar{\sigma}$, step size h , SubTB length-mixing λ , and trajectory-sampling strategy materially affect stability/convergence.

Personal Opinion on this work

This paper addresses a clear problem with path-space samplers, which is that both PIS and DDS required full trajectory signals to learn. By introducing subtrajectory signals from the GFlowNet framework, they were definitely able to reduce the credit assignment problem. Before being able to use any of these path samplers in settings where we only have the unnormalized density function, mainly in chemistry and physics, we will need much better mode coverage since these unnormalized densities are highly multidimensional. The difficult benchmarks have shown that we are not there yet, but there is still a lot of work to do on mode discovery and sample efficiency in order to use these models in real-life applications.

Future Work

Rare Event Sampling. Extend DGFS to efficiently sample from rare event distributions, where the target density is concentrated in low-probability regions of the reference process.

Movement from Path Space to Latent Space. Many diffusion models operate in latent spaces for high-dimensional data (e.g., images). Adapting DGFS to learn flows in latent spaces could improve scalability and efficiency. It means they don't have to start with \mathbb{R}^d as state space but could work in any latent space.





Introduce trick from RL. Incorporate techniques from reinforcement learning, such as experience replay with buffer, off-policy training to improve exploration and sample efficiency in DGFS.

could add something with the partition function estimation Z




Questions and hopefully answers :)

Questions?

References I

-  Bengio, Emmanuel, Salem Bengio, et al. (2023). “GFlowNet Foundations”. In: *Journal of Machine Learning Research* 24, pp. 1–55.
-  Bengio, Emmanuel, Moksh Jain, et al. (2021). “Flow Network based Generative Models for Non-Iterative Diverse Candidate Generation”. In: *Advances in Neural Information Processing Systems* 34, pp. 27381–27394.
-  Chen, Nan et al. (2022). “On the convergence of score-based generative modeling”. In: *arXiv preprint arXiv:2206.04114*.
-  De Bortoli, Valentin et al. (2022). “Convergence of score-based generative modeling for general data distributions”. In: *Advances in Neural Information Processing Systems* 35, pp. 21892–21904.

References II

-  Lee, Hyeontae et al. (2022). “Convergence of denoising diffusion models under the manifold hypothesis”. In: *arXiv preprint arXiv:2205.12551*.
-  Zhang, Dinghuai et al. (2022a). “Path Integral Sampler: Diffusion-based Sampling for Unnormalized Densities”. In: *Advances in Neural Information Processing Systems 35*, pp. 11738–11751.
-  — (2022b). “Path Integral Sampler: Diffusion-based Sampling for Unnormalized Densities”. In: *Advances in Neural Information Processing Systems 35*, pp. 11738–11751.